

Loan Approval Prediction

1. Introduction

Loan approval is a critical process in the financial sector. Predicting whether a loan applicant is eligible for approval based on key attributes can help banks and financial institutions make data-driven decisions. This project uses machine learning to build a predictive model that determines loan eligibility based on applicant details.

2. Problem Statement

Financial institutions face risks when approving loans, and manually evaluating applications is time-consuming. This project aims to develop a machine learning model to automate loan approval decisions using historical loan application data.

3. Data Description

The dataset consists of various attributes related to loan applicants, including:

- **Loan_ID:** Unique identifier for each loan
- **Gender:** Male or Female
- **Married:** Marital status
- **Dependents:** Number of dependents
- **Education:** Graduate or Not Graduate
- **Self_Employed:** Applicant is self-employed or not
- **ApplicantIncome:** Applicant's income
- **CoapplicantIncome:** Co-applicant's income
- **LoanAmount:** Loan amount applied for
- **Loan_Amount_Term:** Term of the loan in months
- **Credit_History:** Credit history of the applicant
- **Property_Area:** Urban, Semi-urban, or Rural
- **Loan_Status:** (Target Variable) Loan approval status (Yes/No)

4. Methodology

1. Data Preprocessing

- Handle missing values
- Convert categorical variables into numerical format
- Normalize numerical features

2. Exploratory Data Analysis (EDA)

- Visualize data distributions
- Identify patterns and correlations

3. Model Selection and Training

- Use machine learning models like Logistic Regression, Decision Tree, and Random Forest
- Train models using training data
- Evaluate performance using accuracy, precision, recall, and F1-score

4. Model Evaluation

- Test the model on unseen data
- Analyze performance metrics

5. Implementation Steps

This code is designed to **predict whether a loan application will be approved or not** using machine learning, specifically a **Logistic Regression model**. Let's break it down step by step.

1. Import Required Libraries

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

Output:

(No output expected—just loading libraries.)

- **pandas (pd)** – Helps in loading and manipulating data.
- **numpy (np)** – Provides numerical operations support.
- **train_test_split** – Splits data into training and testing sets.
- **StandardScaler** – Standardizes (normalizes) data for better model performance.
- **LogisticRegression** – A machine learning algorithm used for classification.
- **accuracy_score** – Evaluates how well our model is performing.

2. Load the Dataset

```
df = pd.read_csv("loan_data.csv")
print(df.head())
```

Output :

```
[annapurnagollapalli@annapurnas-MacBook-Air ~ % cd Documents
annapurnagollapalli@annapurnas-MacBook-Air Documents % python3 loan_prediction.py
First 5 rows of the dataset:
   Loan_ID Gender Married ... Credit_History Property_Area Loan_Status
0  LP001002    Male     No ...           1.0      Urban          Y
1  LP001003    Male    Yes ...           1.0     Rural          N
2  LP001005    Male    Yes ...           1.0      Urban          Y
3  LP001006    Male    Yes ...           1.0      Urban          Y
4  LP001008    Male     No ...           1.0      Urban          Y

[5 rows x 13 columns]

Dataset Information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Loan_ID           614 non-null    object  
 1   Gender            601 non-null    object  
 2   Married           611 non-null    object  
 3   Dependents        599 non-null    object  
 4   Education         614 non-null    object  
 5   Self_Employed     582 non-null    object  
 6   ApplicantIncome   614 non-null    int64   
 7   CoapplicantIncome 614 non-null    float64 
 8   LoanAmount        592 non-null    float64 
 9   Loan_Amount_Term  600 non-null    float64 
 10  Credit_History    564 non-null    float64 
 11  Property_Area     614 non-null    object  
 12  Loan_Status       614 non-null    object  
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB
None
```

- Reads the **loan_data.csv** file and stores it in a variable df (DataFrame).
- This dataset contains details about applicants like **income, loan amount, credit history, etc.**

3. Handle Missing Values Properly

In real-world data, some values might be **missing** (NaN). We fill these missing values using appropriate strategies.

```
df['Gender'] = df['Gender'].fillna(df['Gender'].mode()[0])
```

```
df['Married'] = df['Married'].fillna(df['Married'].mode()[0])
```

```
df['Dependents'] = df['Dependents'].replace('3+', 3).fillna(0).astype(float) # Fix NaNs and convert to float
```

```
df['Self_Employed'] = df['Self_Employed'].fillna(df['Self_Employed'].mode()[0])
```

```
df['LoanAmount'] = df['LoanAmount'].fillna(df['LoanAmount'].median())
```

```
df['Loan_Amount_Term'] = df['Loan_Amount_Term'].fillna(df['Loan_Amount_Term'].median())
```

```
df['Credit_History'] = df['Credit_History'].fillna(df['Credit_History'].mode()[0])
```

Output:

```

Missing Values After Cleaning:
Loan_ID          0
Gender           0
Married          0
Dependents       0
Education         0
Self_Employed    0
ApplicantIncome   0
CoapplicantIncome 0
LoanAmount        0
Loan_Amount_Term 0
Credit_History    0
Property_Area     0
Loan_Status        0
dtype: int64

```

- **fillna()** replaces missing values:
 - For **categorical columns** (Gender, Married, etc.), we use **mode** (most frequent value).
 - For **numerical columns** (LoanAmount, Loan_Amount_Term), we use **median** (middle value).
- **replace('3+', 3)** converts "3+" in Dependents column into 3.
- **astype(float)** ensures numerical consistency.

4. Convert Categorical Variables into Numerical (One-Hot Encoding)

```
df = pd.get_dummies(df, columns=['Gender', 'Married', 'Self_Employed', 'Education', 'Property_Area'], drop_first=True)
```

Output :

```

Dataset after Encoding:
   Loan_ID  Gender  Married  ...  Credit_History  Property_Area  Loan_Status
0  LP001002      1       0  ...        1.0            2            1
1  LP001003      1       1  ...        1.0            0            0
2  LP001005      1       1  ...        1.0            2            1
3  LP001006      1       1  ...        1.0            2            1
4  LP001008      1       0  ...        1.0            2            1
[5 rows x 13 columns]

```

- Machine learning models **don't understand text**, so we convert categorical columns into numerical using **One-Hot Encoding**.
- Converts categorical variables (e.g., Gender, Married) into numeric values (0s and 1s) for machine learning.
- `drop_first=True` removes redundant columns to avoid duplication.

5. Define Features (X) and Target Variable (y)

```
X = df.drop(columns=['Loan_ID', 'Loan_Status'])
```

```
y = df['Loan_Status'].map({'Y': 1, 'N': 0}) # Convert Loan_Status to numerical
```

- X contains **all features** (applicant details) except **Loan_ID** (not useful for prediction) and **Loan_Status** (target variable).
- Defines **X (features)** and **y (target variable)**.

- y is the **target variable**, where:
 - "Y" (Loan Approved) $\rightarrow 1$
 - "N" (Loan Rejected) $\rightarrow 0$

Output:

```
(614, 12) (614, )
```

(614 rows, 11 features)

6. Check for Remaining Missing Values

```
print("Missing values in dataset:\n", X.isnull().sum().sum()) # Should print 0
```

Output:

```
Missing values in dataset:  
0
```

- Ensures there are **no missing values** in X before proceeding.

7. Split Data into Training and Testing Sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Output:

```
(491, 12) (123, 12)
```

- `train_test_split()` **divides data into training and testing sets**:
 - **80%** for training ($X_{\text{train}}, y_{\text{train}}$)
 - **20%** for testing ($X_{\text{test}}, y_{\text{test}}$)
- `random_state=42` ensures we get the same split every time for consistency.

8. Ensure No Missing Values in Train & Test Sets

```
X_train.fillna(0, inplace=True)
```

```
X_test.fillna(0, inplace=True)
```

Output:

(No direct output)

- If there are any missing values left, we fill them with **0** to avoid errors.

9. Standardize the Data (Feature Scaling)

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

Output:

(No direct output—data is now scaled.)

- Standardization **helps models perform better** by making sure all features are on a similar scale.
- **Why?** Features like income (big numbers) can dominate over small-scale features like Credit History.

10. Train Logistic Regression Model

```
model = LogisticRegression()
```

```
model.fit(X_train_scaled, y_train)
```

Output:

(No direct output—model is now trained.)

- Creates a **Logistic Regression model** and trains it using the **training data**.

11. Make Predictions & Calculate Accuracy

```
y_pred = model.predict(X_test_scaled)
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print(f"Model Accuracy: {accuracy:.4f}")
```

Output:

```
Model Accuracy: 0.7886
```

- `model.predict(X_test_scaled)` makes predictions on test data.
- `accuracy_score(y_test, y_pred)` calculates how many predictions were correct.

12. Test Model on a Single Example

```
sample_input = X_test.iloc[0:1] # Take one sample from test set
```

```
sample_input_scaled = scaler.transform(sample_input) # Scale the input
```

```
prediction = model.predict(sample_input_scaled)
```

```
print(f"Loan Approval Prediction: {'Approved' if prediction[0] == 1 else 'Not Approved'}")
```

Output:

```
Loan Approval Prediction: Approved
```

- **Takes one test case**, scales it, and makes a prediction.
- Prints "Approved" if the model predicts **1**, otherwise prints "Not Approved".

Final Complete Code :

```

import pandas as pd

import numpy as np

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import accuracy_score

# Load dataset

df = pd.read_csv("loan_data.csv")

# Handle missing values properly

df['Gender'] = df['Gender'].fillna(df['Gender'].mode()[0])

df['Married'] = df['Married'].fillna(df['Married'].mode()[0])

df['Dependents'] = df['Dependents'].replace('3+', 3).fillna(0).astype(float) # Fix NaNs and
convert to float

df['Self_Employed'] = df['Self_Employed'].fillna(df['Self_Employed'].mode()[0])

df['LoanAmount'] = df['LoanAmount'].fillna(df['LoanAmount'].median())

df['Loan_Amount_Term'] = df['Loan_Amount_Term'].fillna(df['Loan_Amount_Term'].median())

df['Credit_History'] = df['Credit_History'].fillna(df['Credit_History'].mode()[0])

# Convert categorical variables into dummy variables

df = pd.get_dummies(df, columns=['Gender', 'Married', 'Self_Employed', 'Education',
'Property_Area'], drop_first=True)

# Define features and target variable

X = df.drop(columns=['Loan_ID', 'Loan_Status'])

y = df['Loan_Status'].map({'Y': 1, 'N': 0}) # Convert Loan_Status to numerical

print(X.shape, y.shape) # Checking dataset size

```

Check for any remaining NaN values before splitting

```
print("Missing values in dataset:\n", X.isnull().sum().sum()) # Should print 0
```

Split the dataset

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
print(X_train.shape, X_test.shape) # Check sizes
```

Ensure no missing values in train and test sets

```
X_train.fillna(0, inplace=True)
```

```
X_test.fillna(0, inplace=True)
```

Standardize features

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

Train Logistic Regression Model

```
model = LogisticRegression()
```

```
model.fit(X_train_scaled, y_train)
```

Predictions and accuracy

```
y_pred = model.predict(X_test_scaled)
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print(f"Model Accuracy: {accuracy:.4f}")
```

Example prediction

```
sample_input = X_test.iloc[0:1] # Take one sample from test set
```

```
sample_input_scaled = scaler.transform(sample_input) # Scale the input
```

```
prediction = model.predict(sample_input_scaled)
```

```
print(f"Loan Approval Prediction: {'Approved' if prediction[0] == 1 else 'Not Approved'}")
```

Final Output :

```
[annapurnagollapalli@annapurnas-MacBook-Air Documents % python3 loan_predict  
ion.py  
(614, 12) (614,)  
Missing values in dataset:  
0  
(491, 12) (123, 12)  
Model Accuracy: 0.7886  
Loan Approval Prediction: Approved
```

6. Results

- **0 missing values** (data is clean ✓)
- **Accuracy:** 78 % (model performs well ✓)
- **Loan Application result** for a test case is "**Approved**" or "**Not Approved**" .

7. Future Enhancements

- Implement deep learning techniques for improved accuracy.
- Deploy the model as a web application for real-time loan approval.
- Use additional financial data to improve prediction accuracy.

8. Conclusion

This project successfully builds a machine learning model for predicting loan approval. The results demonstrate its potential to assist financial institutions in making data-driven decisions efficiently.